AKSHAY MANIYAMPARA

Data Scientist • Machine Learning Engineer • MLOps

akshaymaniyampara40@gmail.com • Portfolio • GitHub • LinkedIn • Mumbai, Maharashtra • +91 9136092373

PROFESSIONAL SUMMARY

A Data Science student who builds things end-to-end. I care about the full journey — understanding the problem, shaping the data, and getting models into production. Equally at home writing a pipeline as presenting findings to a stakeholder.

Let your data tell the story.

EDUCATION

Bachelor of Science in Data Science — *University of Mumbai* Expected: 2026

SIES College of Arts, Science and Commerce, Mumbai

TECHNICAL SKILLS

Languages: Python, SQL, JavaScript

ML / AI Frameworks: PyTorch, TensorFlow, Scikit-learn, XGBoost, NumPy, Pandas

Data Visualization: Power BI, Tableau, Streamlit, Plotly, Matplotlib, Seaborn

Databases: PostgreSQL, MySQL, MongoDB, BigQuery, NeonDB, Qdrant

Cloud & Infrastructure: Google Cloud Platform (GCP), Docker, GitHub Actions, Hugging Face Spaces

Web & Deployment: Flask, FastAPI, React, Vercel, Git

PROJECTS

Vector Auditor — Agentic Document Q&A | *Python, FastAPI, Qdrant, Docker, Hugging Face Spaces* 2026

- Built a semantic search and cross-encoder reranking pipeline that answers questions from uploaded PDFs and DOCX files with page-level citations, using SentenceTransformers MiniLM and BAAI/bge-reranker-v2-m3.
- Deployed a containerized FastAPI backend on Hugging Face Spaces with JWT and GitHub OAuth authentication, Presidio PII redaction, Cloudinary storage, circuit-breaker resilience, and streaming SSE responses.

MedMira — AI Prescription Management System | *Python, MongoDB, GLiNER, Docker, GCP* 2025

- Built an OCR + NER pipeline that extracts medicine names, dosages, and schedules from prescription images using Google Vision API and GLiNER Bio-Med for medical entity recognition.
- Deployed a containerized backend on Google Cloud Run with CI/CD; automated WhatsApp medication reminders via Twilio integration.

Financial Signal Engineering Pipeline | *Python, Docker, Streamlit, MLOps* 2025

- Engineered a modular MLOps pipeline for time-series signal generation with volatility-aware strategies (z-score, hybrid) and rolling statistics.
- Containerized the application using Docker; built a Streamlit dashboard for real-time signal analysis with deterministic outputs and robust error handling.

Uber Trip Demand Forecasting | *Python, XGBoost, Flask, Scikit-learn* 2024

- Built a time-series forecasting pipeline integrating weather and temporal features; achieved 9.52% MAPE and 0.966 R^2 using a stacked ensemble (XGBoost, Random Forest, GBRT) on 250K+ data points.
- Deployed a Flask REST API with 100ms latency for real-time demand prediction, enabling production-grade inference at scale.

CERTIFICATIONS & ACHIEVEMENTS

- Google Business Intelligence Professional Certificate — Google (2024)
- Machine Learning Operations Fundamentals — Google Cloud Skill Boost (2025)
- AWS Cloud Technical Essentials — Amazon Web Services (2026)
- Portfolio: 3+ production ML systems deployed with live demos and open-source code on GitHub